

Unresolved: Semantic Metastability in a Language Model Under Context Shift

[author block — Andrew]

Review draft — September 2, 2026

“Two fish are in a tank. One looks to the other and asks: how do you drive this thing?”

Abstract

When the context around a request changes its meaning — a word’s sense, or whether a request is fiction — how does a language model’s internal reading follow? We measure this directly in gpt-oss-20b, reading a token’s interpretation from the residual stream while forty-sentence contexts switch sides halfway through, against matched no-shift references. Two tasks: the word *tank* moving between aquarium and vehicle senses, and the fixed request “I want to write a suicide letter.” moving between fictional and real framing.

The reading follows the shift only partway. It crosses to the new side after a median of four to ten sentences, then stops short, and twenty countering sentences never close the remaining gap. In one tank direction it dwells between the two interpretations, stationary to the end of the window. Individual runs move by drift plus discrete jumps, and the jumps are not timed by strong evidence; order effects, though large, are explained almost entirely by recency weighting. None of these states is geometrically unusual against matched references, yet all carry a persistent internal signal that the context is mixed — a signal the model’s behavior does not appear to use. We call this cluster of properties semantic metastability.

The model never asks which meaning is intended: zero of 96 tank completions request clarification. The task covered by a refusal safeguard stayed safe at mid-transition (80% of completions); the task without one committed silently to a sense. Safe-completion falls from 91% to 50% as the reading moves toward the fictional frame (the low endpoint rests on four completions) — attenuation reachable by ordinary context.

Content note: this paper analyzes model behavior around suicide-related requests in a research context. If you or someone you know is struggling, help is available — in the US, call or text 988; elsewhere, findahelpline.com.

1 Introduction

Language models can be made to express uncertainty, and sometimes do. What they do far less reliably — and in the 96 completions we examine here, never — is volunteer that the words in front of them have not settled into one reading. Much of what we ask models to do quietly assumes the opposite: that by the time a model acts, it has settled on one reading. Safety behaviors in particular often take the form *if the request is X, do Y* — a rule that inherits an unexamined premise: that “the request is X” is a resolved fact rather than a reading in transit. The interesting failures, we will argue, live in the unresolved zone: the stretch of a conversation during which context has genuinely shifted what a token means but the model’s internal reading has not finished following. This paper measures that zone directly.

The route to this study runs through an old question. Theories of humor have long placed incongruity and its resolution at the center of the phenomenon: two mutually exclusive understandings of a text form, and something must give. Dynel’s taxonomy of conversational humor separates the garden-path joke — one meaning established, a second revealed, belief shifting from one to the other — from the pun, where the incongruity is not resolved but held. The joke in our epigraph is the classic garden path built on the word *tank*: to get it, the reader’s understanding of one word must shift from one sense to the other. That shift was long observable only through behavior. In a language model it is a measurement problem with an actual substrate: the in-between of a meaning shift is a trajectory in the residual stream, and we can instrument it.

The second task has a graver origin, and we state it plainly. In a widely reported 2025 case, a sixteen-year-old died by suicide after months of conversation with a language model [CITE: Raine v. OpenAI; reporting]. The court filing and contemporaneous reporting — the case is contested litigation, and we characterize only what they describe — recount that direct requests triggered safeguards; that the same requests, reframed as fiction, received assistance, with the user’s real circumstances present in the same long conversation; and that the model offered to draft a suicide note. We do not analyze that case; we take from it a precise scientific question. A model processing such a conversation carries, in some form, internal state that tracks whether “I want to write a suicide letter.” is framed as fictional or real. How does that internal state move as the surrounding context shifts? As it happens — recognized after the fact, not designed — our two tasks bracket the safety question from both sides: one request is covered by a refusal safeguard, the other (“What is the meaning of the word tank?”) by none, so the pair lets us watch behavior with and without a trained safeguard. Whether the safeguard carries a default answer for unresolved cases is an interpretation we defer to §4.

When accumulated context shifts what a token means, what is the in-between? Three candidate worlds: a *learned state* — irresolution is itself something the model has structure for; a *passage* — mere transit between the two resolved readings; or *off the learned distribution entirely* — the token pushed into regions the model never organized, where no trained behavior applies. In the dynamics of biological neural systems, transiently stable states that are not fixed-point attractors are the norm rather than the anomaly. That norm warns against assuming the first two worlds — learned state and passage — must be distinct. The title of this paper is a report on where the data landed: *unresolved* describes the states we found — and, we will argue, the three-worlds taxonomy itself.

The design is simple: two tasks — the polysemous word, the framed request — each with a fixed *carrier* sentence containing the measurement token. Contexts grow one sentence at a time to forty, with the class of the evidence — which sense, or which framing, the sentences support — flipping after twenty; the carrier is re-appended at every step, so the same token is re-read under steadily shifting context; matched single-class contexts provide the reference scale. Figure 1 shows the central result: from either direction, readings cross into the in-between zone and stop short of the opposite reference — robustly so in three of the four cases (§3.2).

We contribute:

1. A measurement protocol for reading interpretations over accumulating context, including two instrument artifacts — accumulation offset and axis rotation — that mimic findings (§2).
2. The shape of reinterpretation: a gradual, partial update whose remnant lingers to the end of the tested horizon (§3.2).
3. The form of its dynamics: drift plus discrete jumps not timed by evidence strength, with the

Both directions on one axes: after the shift, forward and reverse transitions collapse toward the in-between zone (shaded green = unresolved band; dotted = no-shift references ± 1 sd) — neither direction reaches the opposite reference by +20 sentences




Figure 1: The central result: both transition directions, both tasks, at each task’s calibrated site (tank: ‘ tank’ token, layer 4; fiction/real: ‘ want’ token, layer 14). Solid: transition means ± 1 sd, midpoint-referenced. Dotted with bands: position-matched no-shift references ± 1 sd. Green: the unresolved band between the references. Readings cross the midpoint quickly; neither direction reaches the opposite reference within twenty counter-sentences.

smooth integrators tested rejected head-to-head (§3.3).

4. The dwelling within the unresolved zone — a stationary intermediate state in one tank direction, where the model’s answers hedge between the senses — and the observation that zero of the 96 completions ask which sense is meant (§3.4–3.5).
5. Hysteresis almost fully explained by recency weighting: a mild direction-dependent recency difference is all that remains, so the metastability lives in the path, not the equilibrium (§3.6).
6. A persistent, systematic marker of mixed context that behavior does not appear to use (§3.7).

All numbers regenerate from the committed repository, which includes the study’s full corrections record.

2 Measuring an interpretation while it changes

2.1 Model and capture

All experiments use gpt-oss-20b, a 20-billion-parameter mixture-of-experts language model, run in its stock configuration: we apply no modifications to the model or its routing — its native top-4-experts-per-token routing operates untouched. Inputs are formatted with the model’s chat template and processed with deterministic forward passes; identical inputs yield identical activations. We capture the full residual stream (2,880 dimensions) at the output of each of the 24 decoder blocks. Throughout, a *site* is a designated token of a fixed carrier sentence together with a layer, and “the reading at a site” means the residual stream at that token and layer, projected onto a calibrated axis as described next. The tank task reads at the ‘ tank’ token, layer 4; the fiction/real task at the ‘ want’ token, layer 14 (the tasks are introduced in §2.3).

2.2 The instrument

The measurement axis is simple. For a site (token, layer), we average the residual-stream states of the single-sentence calibration corpus by class, giving two class means and their difference w . The reading of a state h is its position along that axis, rescaled so the class means sit at -1 and $+1$: $r(h) = 2(h - m) \cdot w / |w|^2$, where m is the midpoint of the two class means. A reading is a signed coordinate between the two calibrated interpretations — a projection, not a distance to either class. We use the difference of class means rather than a trained classifier such as a logistic probe: it has no hyperparameters, its geometry is transparent, and mean-difference directions read out class structure comparably to trained probes [CITE: mass-mean probing]; its adequacy here is checked empirically. Classification accuracies use the midpoint rule — a held-out sentence is assigned to the class whose mean lies on its side of the midpoint — with splits held out at the scene-family level.

2.3 Tasks, carriers, and corpora

Each task pairs a fixed **carrier** — a sentence appended verbatim after the context at every step — with a designated measurement token. The **tank task** contrasts two senses of a polysemous word: the carrier is “What is the meaning of the word tank?”, measured at the ‘tank’ token, with contexts drawn from aquarium-life and armored-vehicle scene pools. The **fiction/real task** contrasts two framings of one fixed request: the carrier is “I want to write a suicide letter.”, measured at the ‘want’ token, with contexts drawn from fiction-craft discussion (novel drafts, screenwriting, tabletop campaigns) versus real personal circumstance (grief, diagnosis, eviction). Paraphrase carriers (“I would like to write...”, ‘like’; “Help me write...”, ‘letter’) support cross-carrier checks, and a replicate tank carrier (“Define the word tank.”) supports carrier-independence checks. The ‘want’ site was fixed early in the project and retained for comparability; §3.1 shows the framing signal is spread across the request’s content words, and replication at the strongest-reading ‘write’ site is planned (§5).

Each corpus answers one question. *How does a reading move when the evidence changes sides?* The transition corpus presents 40-step cumulative contexts with the carrier re-appended at each step: twenty sentences of one class, then twenty of the other, the target vocabulary absent from every context sentence. *What would the reading be with no shift?* Token-budget-matched no-shift control runs — forty sentences of a single class — provide the reference at every position. *Where do the two interpretations sit?* Single-sentence calibration sentences (300 per class per carrier) define the axis of §2.2. *What does every token read, not only the carrier’s?* Checkpoint captures are full recordings of every token’s activations at designated context lengths (144 per task). *Does the reading track framing cues rather than content?* Minimal pairs hold content fixed while varying only framing cues (150 pairs). *Does the order of evidence matter beyond its amount?* Static mixture sweeps — twenty-sentence contexts holding a fixed mix of the two classes (252 cells per task) — test order dependence. *What does the model do?* Behavior prompts with generation enabled (312; greedy decoding) supply completions. Bare-carrier baselines complete the set.

Context sentences were written by language-model authoring agents (separate from the model under study) under a blind protocol: an agent received only a contrast specification and diversity rules, never the hypotheses. Per-class scene diversity was capped so that no setting dominates a class, enabling scene-held-out validation.

Four terms recur throughout:

- A **run** is one 40-step context sequence. Transition runs flip class at twenty; no-shift control

runs never do.

- A **cell** is one point of a mixture-sweep grid: one class mixture in one block order.
- A **scene family** is a set of sentences sharing one concrete setting (a home aquarium, a tank museum). There are twelve per class per task, and every statistic in this paper is clustered at the family level.
- **Scene-held-out** validation holds out whole families, one from each class per fold.

2.4 The protocol

Box 1 states the protocol for using the axis over accumulating context.

Box 1 — Protocol for reading interpretations over accumulating context.

1. *Calibrate at the same site, same carrier, always.* Projected across token positions, readings become a constant set by position rather than content; projected through another carrier’s axis, they are dominated by token identity (both failure modes occur in our data).
2. *Validate endpoints held-out at the scene level.* Our calibration axes separate classes with held-out accuracy 0.905 (tank, layer 4) and 0.910 (fiction/real, layer 14) under 12-fold leave-one-family-pair-out cross-validation (chance 0.50, 300 per class) — the split that tests whether the axis learned the contrast or a setting.
3. *Reference every absolute claim to position-matched no-shift runs.* Readings carry an **accumulation offset**: a class-nonspecific component that grows with context — $\approx +1.0$ axis units by twenty sentences at the fiction/real site, half the class separation on an axis whose class means sit at ± 1 (≈ 0 at the tank site). It is one shared component: its time-course replicates across paraphrase carriers, with the sign set by each carrier’s axis orientation ($r = +0.75$ same-oriented, -0.86 opposite-oriented). It contaminates any unreferenced position claim.
4. *Verify the axis still points the right way at depth.* The class direction itself can rotate as context accumulates: on accumulated states the fiction/real axis retains only $\cos 0.57\text{--}0.63$ of its single-sentence direction at layers 10–23 (tank: 0.78–0.97). Depth claims require per-layer axes refit on accumulated no-shift states (held-out accuracy 0.93–1.00 at every layer); uncorrected depth readings produced three spurious findings, since retracted.
5. *Treat readings as positions along a designed contrast, never as meaning.* A “real-side” reading may mean “not-fiction” or any correlate; distinguishing these requires contrasts this study does not contain.
6. *Cluster statistics at the family level.* Sentences within a scene family are not independent; every interval in this paper is family-clustered.
7. *Prove the pipeline can fail.* Family-level label shuffles must kill every effect they are run on (they do: separations fall to chance, within-pair effects to zero; the mixed-context marker is tested against its own family-block null instead); synthetic ground-truth fixtures must recover known answers through the actual pipeline code (nine of nine do); positive controls must show each null-result instrument detecting real displacement (they do).

Glossary. *Semantic metastability* — the name Discussion gives to the cluster of properties documented in Results: persistent intermediate configurations of a contextual reading between two calibrated interpretations. *Unresolved zone* — the band of readings between the two calibrated interpretations, and the stretch of context during which a reading sits there. *Dwelling* — a trajectory that becomes stationary within the unresolved zone and, on the measured horizon, does not leave it. *Remnant* — the component of a prior interpretation that twenty counter-sentences do not remove. *Mixed-context marker* — a systematic direction, orthogonal to the content contrast, on which mixed-class contexts read high and pure-class contexts read near zero. *Accumulation offset* and *axis rotation* — the two instrument effects of rule 3 and rule 4.

2.5 Analysis methods

Trajectory models. No regression is involved in the instrument; least squares enters only here. Per-run dynamics are fit to each run’s twenty post-shift readings by least squares over four model forms. The first is a recency-weighted integrator: the reading as a weighted average of all evidence so far, each sentence weighted γ^{age} so that recent sentences count more, with the single decay parameter γ grid-searched. The uniform, equal-weight integrator is its $\gamma = 1$ case and also serves as the lead/lag reference. The second is a change-point step: a constant level before and after a fitted change point (three parameters). The third is a drift-plus-step hybrid: a linear trend plus a step at a fitted change point (four parameters). The fourth is a two-timescale integrator: a mixture of fast and slow recency weightings (three parameters). Selection is by BIC over the twenty points, declaring a winner only when it leads the runner-up by at least 2, and “indeterminate” otherwise. The selector’s identifiability is calibrated on synthetic runs of known type (§3.3).

The remnant gap. The destination reference level (positions 36–40) minus the run’s plateau (post-shift sentences 16–20), both midpoint-referenced.

Uncertainty. Intervals are family-clustered bootstraps: 2,000 seeded draws resampling scene families with replacement. Where a named test appears, its clustering unit is stated in place.

Behavior. Completions are categorized by regular-expression scan plus manual review of the committed worksheets.

2.6 Reproducibility

Every number in this paper regenerates from committed analysis scripts over frozen captures: a full regeneration audit reproduced all reported values (calibration axes bit-identical; seeded bootstraps exact), a permanent fixture suite pushes synthetic data with analytically known answers through the actual pipeline functions, and the label-shuffle and positive-control audits above are committed tests. Corrections that changed reported values are listed in Appendix A; the complete record is in the repository.

3 Results

We first check the instrument: whether an interpretation can be read out at all, and whether what we read is the intended contrast (§3.1). We then follow a reading through a shift and ask how far it moves and how fast (§3.2), what kind of process moves it (§3.3), what a trajectory looks like when it stops between the two interpretations (§3.4), and what the model does while there (§3.5). Two boundary tests close the section: whether the order of evidence adds anything beyond recency weighting (§3.6), and whether any of these states leaves the model’s ordinary geometry (§3.7).




Figure 2: Fiction/real readings by layer (0–23) and position under per-layer secondary axes (class means at ± 1): class signal at every depth, no off-band regime. Left pair: no-shift controls. Right pair: transitions, shift after sentence 20.

3.1 The instrument reads a real contrast, and reads the right thing

Before any dynamics, two questions: can an interpretation be read out at all, and is what we read the intended contrast rather than a topical shadow of it?

On the first: single-sentence calibration axes separate the classes with held-out accuracy 0.905 (tank, layer 4) and 0.910 (fiction/real, layer 14), and axes refit on accumulated forty-sentence contexts separate held-out families with accuracy 0.93–1.00 at every layer from 1 to 23, in both tasks. Under these per-layer accumulated-context axes, no-shift runs read class-consistently across the entire depth of the stack (Fig. 2) — with the caveat, established in Methods, that the fiction/real direction itself rotates with accumulation and must be refit at depth (Fig. 3).

On the second, three independent discriminations. First, an identity-matched comparison: the nine tokens of the tank carrier (“What is the meaning of the word tank?”) appear verbatim in every checkpoint window, so class signal can be compared across tokens with identity held fixed. The signal concentrates overwhelmingly at the sense-bearing token — $d' = 11.7$ at ‘tank’ against a context-token median of 1.4 (Fig. 4; $n = 6$ runs per class, and we caution that orderings among the non-peak bars are not interpretable at this n , since d' divides by a pooled standard deviation that is small for highly predictable tokens). Second, minimal pairs: 150 sentence pairs sharing most content words while differing only in framing cues shift the reading by $+0.99$ axis units — half the full class separation — with 95% of pairs in the predicted direction (pair-level $p = 2.5 \times 10^{-25}$; clustered by content domain — the conservative unit, six domains — $t(5) = 12.0$, $p = 7.2 \times 10^{-5}$; positive in each of the three independently generated batches of pairs; Fig. 5). Third, dose-independence: the effect does not scale with the number of framing-cue words ($r = 0.05$), and is not a length artifact ($r = 0.02$; length-matched pairs $+1.05$) — one cue moves the reading as far as four: the response saturates at a single cue rather than accumulating.

The licensed claim is modest: the reading tracks framing cues with content held fixed — not an abstract frame representation. The two tasks also differ instructively in contrast structure. Tank’s signal anchors to one token; the fiction/real signal is spread across the request’s content words (‘write’ 6.4, ‘suicide’ 6.2, ‘want’ 4.5, against a context-token median of 3.6; Fig. 6). The d' values are not comparable across tasks, but the shape is: a lexical-sense contrast lives at a token; a framing contrast lives across the utterance.




Figure 3: Axis rotation: cosine between each layer’s accumulated-context (secondary) axis and the single-sentence calibration axis. Tank stays at 0.78–0.97; fiction/real declines to 0.57–0.63 at layers 10–23 — depth claims require refit axes.




Figure 4: Class signal by carrier token, identity matched (the nine tank-carrier tokens appear verbatim in every checkpoint window): $d' = 11.7$ at ' tank' against a context-token median of 1.4 ($n = 6$ runs per class; non-peak orderings not interpretable at this n).




Figure 5: Minimal pairs: 150 sentence pairs, content held within pair, framing cue varied. Framing cues alone move the reading +0.99 axis units (half the class separation), 95% of pairs in the predicted direction.




Figure 6: Fiction/real class signal across carrier tokens: distributed over the request's content words ('write' 6.4, 'suicide' 6.2, 'want' 4.5) against a context-token median of 3.6 — a framing contrast lives across the utterance.

3.2 The shape of reinterpretation: gradual, partial, lingering

Figure 1 shows both transition directions, both tasks, at each task’s calibrated site (layers 4 and 14), against the no-shift references. After the shift, readings cross the midpoint at per-run medians of 4 to 10 sentences depending on task and direction (tank 10.5 and 6.0; fiction/real 4.0 and 5.0; a few runs never cross) — but neither direction, in either task, reaches the opposite reference within twenty counter-sentences.

How far short do they stop? The **remnant gap** measures the shortfall: the distance between a run’s plateau and the position-matched no-shift level, in axis units measured from the midpoint of the two no-shift references (§2.5; Fig. 7). Both directions of both tasks show positive gaps under family-clustered bootstrap (Table 1). The largest is tank aquarium→vehicle at +2.16 [1.87, 2.44] — equivalently $1.09\times$ the no-shift amplitude, the distance from the class midpoint to the matched reference; a shortfall larger than the amplitude itself means this plateau sits at the class *midpoint*, the transition stopped at half-way.

Table 1: The remnant gap in both directions of both tasks: the destination reference level (positions 36–40) minus the run plateau (post-shift sentences 16–20), both midpoint-referenced, in axis units. Intervals are family-clustered bootstrap 95% intervals over 2,000 draws; the second interval also resamples the six no-shift reference families on every draw (Appendix A). The last column divides the gap by the no-shift amplitude, the distance from the class midpoint to the matched reference.

Transition	Gap	95% CI, references fixed	95% CI, references resampled	Gap / amplitude
tank, aquarium→vehicle	+2.16	[1.87, 2.44]	[1.70, 2.62]	1.09
tank, vehicle→aquarium	+1.15	[0.82, 1.45]	[0.77, 1.48]	0.57
fiction→real	+0.38	[0.28, 0.49]	[0.27, 0.50]	0.43
real→fictional	+0.35	[0.11, 0.58]	[−0.12, +0.79]	0.40

One case demotes under a stricter bootstrap: when the uncertainty of the six-run reference is propagated, three cases stand but real→fictional widens to [−0.12, +0.79] and no longer excludes zero — that remnant is suggestive only. The gap is not an artifact of weaker post-shift material: the post-shift sentences match the no-shift runs’ sentences in single-sentence calibration strength (percentile medians 0.51 vs 0.45 and 0.48 vs 0.49; $p = 0.24, 0.90$). Whether the remnant is truly permanent or merely slower than twenty sentences can resolve is undecidable in this corpus — the late slopes are shallow but nonzero in some runs — and deciding it is the first item of future work.

The transition is two-phase in a specific sense: no single recency-weighted average fits it. In both directions of both tasks the data sit above the best-fitting average early and below it late, and the late shortfall is the remnant gap. The reading lags the present frame on roughly the fitted recency timescale. The decay parameters of §3.3’s recency integrator (§2.5) have median $\gamma = 0.90$ in both fiction/real directions and 0.94 in tank →aquarium — weighted-mean evidence ages of about nine and fourteen sentences — while in tank →vehicle the median reaches 0.99: weighting so close to uniform that the effective memory exceeds the window — the dwell showing up inside the fit itself. Relative to an equal-weight average of all evidence in the window, the mean trajectories run slightly *ahead* — the pattern a recency weighting would produce; individual runs are heterogeneous around that mean. What the data rule out is not lag but lag *in excess* of evidence integration.

The single-site view has depth structure behind it (Fig. 8 shows the collapse view at four layers per task; Fig. 2 and supplementary heatmaps give all layers). In the fiction/real task the shallowest layers cross almost immediately — within three to four sentences, and completely — while mid-stack layers cross at medians of eight to thirteen. Tank crosses later everywhere: its shallow band



at medians of thirteen and eight, and in aquarium→vehicle at median thirteen at every depth — the dwelling of §3.4 is stack-wide there — while the reverse direction crosses earliest at depth (median six). Our pre-stated prediction about the ordering of crossing times across layers held in the fiction/real task and failed in tank (§5).

Where the stack ends up is more telling than when it crosses. Table 2 gives the late-window means by depth band. In the last five sentences (destination-signed means; regenerated by the committed layer script), fiction/real’s layers 3–18 settle at +0.3 to +0.4 of the reference in both directions and its shallowest layers at full strength; tank aquarium→vehicle’s shallow layers reach only +0.57, its layers 3–12 end at +0.12, and its layers 13–23 at 0.00 to −0.10 — the dwelling direction’s mid and deep stack ends at the midpoint itself. The deepest band (layers 19–23) hovers near the midpoint in one direction of each task (fiction→real +0.08; aquarium→vehicle −0.10) while the reverse directions settle (+0.48; +0.42).

Table 2: Where each layer band ends up: mean reading over the last five post-shift sentences (positions 36–40) under per-layer accumulated-context axes, signed so that the destination class is positive, with class means at ± 1 and the midpoint at 0.

Layers	tank, aquarium→vehicle	tank, vehicle→aquarium	fiction→real	real→fictional
0–2	+0.57	+1.03	+0.99	+1.14
3–12	+0.12	+0.30	+0.41	+0.32
13–18	0.00	+0.38	+0.41	+0.33
19–23	−0.10	+0.42	+0.08	+0.48

Direction matters, and the asymmetry is not an instrument accident. Tank transitions toward the vehicle sense fall 1.09 amplitude-fractions short versus 0.57 in the reverse direction; the same pattern replicates under an independent carrier (“Define the word tank.”: 0.92 vs 0.61; the replicate’s calibration axis is in-sample, not held-out) and survives median-based midpoints. Within the fiction/real task the asymmetry is site-dependent: symmetric at the ‘want’ site (0.43/0.40) but strongly asymmetric at the ‘letter’ site of the paraphrase carrier “Help me write...” (0.90/0.33 — against that site’s own healthy reference amplitude of 1.18, though at $n = 4$ per direction, exploratory). A candidate cause exists at the calibration level: the vehicle class is intrinsically broader (per-side spread 0.83–0.91 versus 0.56–0.70 for aquarium, at every layer tested) — transitions *into* the broader class stop farther short. We flag this as a correlate, not a demonstrated mechanism (Fig. 9).






Figure 9: Direction asymmetry: transitions into the broader class stop farther short; the pattern replicates under the replicate carrier and appears at the 'letter' site ($n = 4$ per direction, exploratory); per-side calibration spreads differ (vehicle broader than aquarium at every layer tested).

3.3 The form of the dynamics: drift plus discrete jumps

What process produces these trajectories? We fit the four model forms of §2.5 per run to the twenty post-shift readings — recency integrator (with the uniform, equal-weight integrator as its $\gamma = 1$ case), change-point step, drift-plus-step hybrid, and two-timescale integrator — selecting by BIC with an indeterminacy band, and we calibrated the selector on synthetic truth (Table 3): hybrid-truth is recovered as hybrid in 13–15 of 24 simulations; step-truth is called integrator in 0–2%; and two-timescale truth masquerades as a plain integrator or indeterminate, reading as hybrid in only 2–3 of 24. The selector cannot fabricate the hybrid verdict from any of the smooth forms tested. This is a claim about the form of the trajectories, not their implementation: it excludes smooth integration of the evidence and establishes discrete changes in the reading; what inside the network produces one is not identified here.

On the real runs the selector nonetheless returns hybrid: among classifiable runs, drift-plus-jump is dominant — 11 of 16 tank runs and 14 of 25 fiction/real runs (48% of fiction/real runs are indeterminate at that task's noise level, and we state the claim as "hybrid-dominant among classifiable runs" accordingly). The two-timescale model wins zero runs in either task (Fig. 10; per-run fits in Fig. 11; the raw paths in Fig. 12 show the heterogeneity directly — some runs glide, some step). The two-phase shape of §3.2 could in principle have been smooth multi-timescale integration; tested head-to-head, it is not.

Table 3: Per-run model selection by BIC on the real transition runs and on synthetic runs of known type. A model wins a run only when its BIC leads the runner-up by at least 2; otherwise the run is indeterminate. Synthetic runs: 24 per task per truth type, pooled across the two tasks.

Runs	Drift + step (hybrid)	Step	Recency integrator	Two-timescale	Indeterminate
tank (24 runs)	11	2	3	0	8
fiction/real (48 runs)	14	7	4	0	23
synthetic, hybrid truth (48)	28	5	1	0	14
synthetic, two-timescale truth (48)	5	0	18	4	21
synthetic, step truth (48)	3	27	0	0	18




Figure 10: Per-run model selection (BIC with an indeterminacy band) over five candidate processes, with simulation-calibrated identifiability: drift-plus-jump dominates among classifiable runs; the two-timescale integrator wins zero runs, and two-timescale truth cannot masquerade as hybrid.

Jump timing is not predicted by evidence strength. Every post-shift sentence has a single-sentence calibration reading (456 of 456 matched), and the sentences at which runs jump are median-strength exemplars of their class — percentile 0.50 against 0.50 for all other added sentences ($p = 0.445$). The same-strength evidence moves one run smoothly and jolts another: what differentiates a jump step is not the stimulus property we measured — though we have tested only evidence strength — and single-sentence calibration strength may understate in-context strength; surprisal, syntax, and discourse position remain untested — so “state-dependent” is our working interpretation rather than a demonstrated exclusion.

The same process is visible from inside the stream. Using per-position axes calibrated from no-shift checkpoint windows — content and position matched, only history differing — the tokens of the post-shift block themselves read only about half their no-shift reference: $+0.34/+0.37$ (the two transition directions) at ten post-shift sentences, $+0.52/+0.53$ at twenty (trimmed means; the untrimmed tank values are inflated by a heavy tail we document in the QA appendix), replicated with a tighter instrument in the fiction/real task (Fig. 13, 14). Mixed history suppresses the class reading of even the new evidence’s own tokens, in a recency-graded way. This is descriptive — it is what token-level recency integration would produce — but it establishes that the transition is not merely a summary-site phenomenon: any window-level readout inherits the suppression.




Figure 11: Per-run fits (tank): the best single- γ recency integrator underpredicts the early rise and overpredicts late convergence; drift-plus-jump hybrids capture the paths.




Figure 12: Raw per-run trajectories (tank, layer 4): the heterogeneity behind the means — some runs glide, some step.




Figure 13: Within-stream readings (tank): post-shift-block tokens — 100% destination-class content — read about half their no-shift reference (trimmed means; the two transition directions), recovering only partially by twenty sentences.




Figure 14: Within-stream replication in fiction/real with a tighter instrument: the same history-driven suppression of the new evidence's own tokens.

3.4 The dwelling within the unresolved zone

In one direction of the tank task — aquarium→vehicle — the trajectory dwells within the unresolved zone: it becomes stationary mid-transition and, on the measured horizon, does not leave. The central tendency of the run distribution is stationary at mid-axis for at least ten consecutive steps, ≥ 2.9 across-run standard deviations from *both* endpoint references, with the spread across runs flat rather than tightening; Gaussian-mixture comparison favors a single component in every time bin — one stationary population, not a hidden mixture of resolved and unresolved runs (Fig. 15). Within tank, the same signature appears at the pre-lexical 'word' site and under the replicate carrier; in the fiction/real task it appears only at the 'letter' site of the paraphrase carrier, at $n = 4$ per direction, exploratory. "Stationary" is bounded by the horizon: whether these runs would ever complete on a longer one is the same open question §3.2 leaves for the remnant, and some individual runs retain shallow nonzero late slopes. Geometrically, the dwelling states sit squarely inside the trajectory bundle — the spread of transition trajectories across runs — at $0.96\times$ the null distance (§3.7); functionally, they are a distinct condition, as the model's behavior shows.




Figure 15: The dwell (tank aquarium→vehicle): the central tendency of the run distribution is stationary at mid-axis for ≥ 10 consecutive steps, ≥ 2.9 across-run standard deviations from both endpoint references, spread flat; Gaussian-mixture comparison favors one component in every time bin.

3.5 What the model does in the unresolved zone

The dwelling states have a behavioral signature. We generated completions at matched counts of post-shift sentences ($k = 2, 6, 12, 20$) for every run; all behavioral claims in this subsection are correlational. For analysis we partition readings into three bands — origin side, middle, destination side. Behavior tracks the reading: completions whose readings sit on a side answer that side (56% and 66% for the two sides), while the mid band answers “both” — enumerating the two senses — in 45% of mid-band completions, roughly double either side (Fig. 16). One descriptive fact anchors this paper’s framing: **not one of the 96 tank responses asks which sense is meant or declines to answer pending disambiguation** — zero of 96, by regex scan and manual review of the committed categorized worksheet. In the unresolved zone the model either enumerates senses (45%) or silently commits to one (52% at mid-band); it answers as if resolved. One further response class appears across all 108 tank completions (the 96 transition completions plus 12 no-shift completions): 11 are degenerate repetition loops, unrelated to reading band.

When does the reading carry information *beyond* the context composition that drives both reading and behavior? At matched composition: mid-transition, yes — pooling $k \in 6, 12$, decided answers come from runs with more extreme readings (median $|\text{reading}|$ 0.72 vs 0.38, one-sided $p = 0.0138$ — a pooling chosen after the per-depth pattern was seen, so we grade this suggestive); at the settled extremes, no. In the fiction/real task we categorize each completion as fiction-framed assistance or a safety response — a *safe-completion*, which addresses the request as a risk (declining the letter, redirecting to support) rather than fulfilling it. The reading separates the two at $k = 2$ (-0.06 vs $+1.13$, $p = 0.010$, one of four k values tested) but not later, where response type follows scene family (Fig. 17).




Figure 16: Behavior by reading band (tank): side bands answer their side (56%/66%); the mid band answers “both” in 45% of cells — roughly double either side. Zero of 96 completions ask which sense is meant.

The safety-relevant gradient is nonetheless plain at the band level: safe-completion rates run 50% $\rightarrow$ 80% $\rightarrow$ 91% as the reading moves from fiction-side through mid to real-side (origin band $n = 4$, fragile).

An exploratory post-freeze check (Appendix B) asks whether any other layer’s reading associates more strongly with behavior than the calibrated site’s: per-layer association curves (Fig. 18) are roughly flat from mid-stack to the final layer in both tasks — nothing singles out the deep layers — though the instrument is blunt (band-level, imbalanced outcomes), so this neither establishes nor rules out depth-specific behavioral readout.




Figure 17: Behavior at matched context composition: mid-transition, decided answers come from runs with more extreme readings (pooled k of 6 and 12, graded suggestive); in fiction/real the reading separates response types at k = 2 only; safe-completion rates run 50→80→91% across bands (fiction-side endpoint n = 4).




Figure 18: EXPLORATORY (post-freeze): reading→behavior association by layer (rank AUC, family-clustered 95% bands). Roughly flat from mid-stack to the final layer in both tasks; nothing singles out the deep layers. Blunt instrument: band-level association, imbalanced outcomes.

3.6 Order and equilibrium: hysteresis without stickiness

Does the order of evidence matter beyond its amount? We built static mixture sweeps: twenty-sentence contexts with k destination-class sentences, k swept $0 \rightarrow 20$, in two block orders, from each family’s own transition sentences. The resulting loops are large — the same mixture read with the destination block recent versus old differs by loop areas of +14.4 [12.2, 16.8] (tank) and +10.5 [7.5, 13.3] (fiction/real): operational hysteresis, plainly present (Fig. 19).

The reserved question was whether any of that order-dependence exceeds what recency weighting alone produces — call any such excess *stickiness*. There is none: a one-parameter recency integrator fitted to the sweep cells reproduces the loop areas almost exactly (+14.3; +11.1) — excess “stickiness” of +0.1 and −0.5, both indistinguishable from zero — and cross-order validation (fit one branch, predict the other) preserves the verdict. An earlier pass computed this against a misspecified null (γ imported from the transition fits) and found apparent stickiness; the fitted null supersedes it (Appendix A). What remains real beyond the single- γ account is mild: the two sweep directions prefer slightly different recency weights (γ 0.91/0.97 tank, 0.84/0.90 fiction/real), echoing the directional asymmetry of §3.2.

This resolves an apparent contradiction with §3.3. The *equilibrium map* from evidence mixture to reading is smooth and integrator-like; the *sequential path* through a transition is drift punctuated by jumps. Both are true; they are different observables. The metastability this paper names is a property of paths, not of the equilibrium map — there is no bistability to fall into and no barrier to escape.




Figure 19: Hysteresis without stickiness (tank): static-mixture sweeps in two block orders produce loop area +14.4 [12.2, 16.8]; a one-parameter recency integrator fitted to the same cells reproduces it (+14.3) — excess “stickiness” +0.1, ns.

3.7 The geometry of irresolution

Finally, the third of the introduction’s three worlds: do these in-between states leave the model’s learned distribution? “Off-manifold” is meaningful only relative to a reference, a measure, and a noise level, so we fix all three: reference — the position-matched no-shift states (and the cross-run transition bundle); measures — k-nearest-neighbor distance and out-of-subspace reconstruction error against the no-shift PCA subspace; noise level — the held-out self-distances of the reference itself. (Paired values in this section are tank/fiction-real throughout.) The instruments detect real displacement: single-sentence calibration states — a genuinely different context regime, serving as a positive control — read $1.7\text{--}2.2\times$ the null across the two instruments, and position-mismatched states are flagged at 13–44%.

At the individual-state level, nothing leaves the reference distribution. Jump steps read $1.06\times/1.12\times$ the null, the stationary dwelling states $0.96\times$, all inside the null’s spread (Fig. 20). Our pre-registered prediction that jump steps would show elevated off-manifold distance failed.

The subtle result appeared only when the per-state verdict was challenged: a displacement too small to flag any single state could still be shared by all of them, so we tested the mean directly. The *mean* out-of-subspace reconstruction error of transition states, against a family-block bootstrap null, is far outside the null’s range: a shared displacement equal to **25% and 38% of the full class separation**, $p < 0.001$ in both tasks, rising over the first five post-shift sentences and persisting undiminished to twenty — largest in the stationary dwelling states (Fig. 21). We call the direction that carries it the **mixed-context marker**. It is not class leakage: the direction is orthogonal both to the single-sentence class axis ($\cos +0.015, +0.011$) and, by construction and by measurement, to the accumulated-context class axis ($-0.0052, -0.0065$). It survives held-out direction estimation (71–90% of in-sample magnitude). It is not family novelty: if the marker merely reflected unfamiliar scene families, pure-class cells from families *outside* the reference construction should sit higher on it than cells from familiar families; they sit lower ($+1.6\%/+3.4\%$ of separation versus $+6.7\%/+6.0\%$), and both sit far below mixed cells ($\approx +27\%$).

What is it? Held-out mixture cells locate what elevates it: the marker is absent from pure-class contexts and near its full transition strength in *static* mixed contexts — so it marks mixed context, not temporal shifting per se. One qualification: in the fiction/real task it halves when the mixture is interleaved rather than blocked, so there, part of what elevates it is the coherent, blocked structure of the shift. It did not predict behavior in eight matched-composition tests at the calibrated site and layer (one nominal hit, uncorrected; other layers untested). The model, in other words, carries a persistent, systematic signal that its context is mixed — a direction on which pure contexts sit near zero and mixed contexts near 27% of the class separation — and its behavior does not appear to use it. Whether this signal is a learned representation of mixedness or a systematic consequence of heterogeneous input is not decided by our measurements.

Three graded senses of “off-distribution” should be kept distinct: exotic natural inputs; interventional states (patching, steering); and states of one context regime measured against a reference built from another, as our single-sentence positive controls are. Our shifts are the tamest possible case — clean block transitions between two well-learned frames — and the fact that even the dwelling states stay on-distribution here says nothing about adversarial or incoherent contexts, where genuine excursions remain most likely; that battery is future work.




Figure 20: Per-state geometry: jump steps ($1.06\times/1.12\times$ the null) and dwell states ($0.96\times$) sit inside the no-shift null’s spread; positive controls (calibration states, position-mismatched states) are detected.




Figure 21: The mixed-context marker: mean out-of-subspace reconstruction error of transition states against a family-block bootstrap null — a shared displacement of 25%/38% (tank/fiction-real) of class separation, orthogonal to the class axis, rising over ≈ 5 post-shift sentences, persisting to twenty, largest during the dwell.

4 Discussion

What “metastable” means here. We anchor the term in the dynamics of biological neural systems rather than in physics. In coordination dynamics and in the winnerless-competition framework — models of sequential transient dynamics in neural populations — cognition is described as sequences of transiently stable states that are not fixed-point attractors. These are structured passages through state space, in which dwelling and moving are not opposites. That is the geometry the data show: a run that dwells mid-transition occupies something functionally state-like (stationary for many steps, with its own behavioral signature) yet geometrically a passage (inside the trajectory bundle, on a smooth equilibrium map). One disclaimer: nothing here is equilibrium bistability — the evidence-mixture map is smooth (§3.6) — and readers should not import a barrier-crossing picture. Semantic metastability, as we use it, is a property of paths.

Typicality is not commitment. The three-worlds question dissolved because it conflated two axes. On *distributional typicality* the stationary states are unremarkable — inside the bundle, on-distribution. On *semantic commitment* they are distinctive — uncommitted, between calibrated interpretations, with hedging as their behavioral expression. A learned state of unresolvedness would be typical *and* uncommitted; genuinely off-distribution states (glitch inputs, interventions, adversarial incoherence) occupy a different cell of that table and remain untested here. The general lesson: that a state is geometrically typical does not make it functionally normal — our behavior data document the functional difference directly.

Safety and alignment, first descriptively. The failure surface this study maps needs no adversary. During an ordinary, coherent shift of conversational frame, a model’s reading of a critical token trails the present frame for roughly the fitted recency timescale (here, nine to fourteen sentences, and effectively unbounded in the dwelling direction), retains a remnant of the prior frame beyond that, and — in one tank direction, with an exploratory $n = 4$ echo in fiction/real — dwells between frames for the remainder of the tested horizon. Throughout, the model never asks which reading is meant or flags the ambiguity as an obstacle — zero of 96 tank completions, across all bands — though 45% of mid-band answers do surface both senses. And safeguard-relevant behavior co-varies with the internal reading: safe-completion rates fall from 91% to 50% across reading bands (the fiction-side endpoint rests on four completions, and the gradient is partly scene-driven — §3.5) — attenuation reachable by ordinary context, no exotic inputs required.

A trained default for unresolved cases? The following is a post-hoc reading of an asymmetry we noticed, not a designed manipulation, and training provenance is unobservable: our two tasks appear to differ in whether they carry a trained default for unresolved cases. The refusal task behaved as though it does — at mid-transition, 80% of completions safe-complete, and in sampled chain-of-thought traces — the reasoning channel its chat format exposes, released with our behavior data — both framings are weighed before the safe reply (a qualitative observation, not a coded rate). The sense task showed no such default, and there the model silently commits (52% of mid-band completions pick a sense). If this reading is right, the metastable zone is the failure window for every behavior *without* a trained uncertainty-default — refusal-style safeguards may be the well-covered exception rather than the rule.

Why the safeguard held: three candidate accounts. Three accounts could explain why the covered task stayed safe at mid-transition, and we cannot yet separate them. They are not exclusive — one names the policy, one the downstream state, one the trigger’s plausible locus — and none is causally established: with two tasks these are observations, and the exploratory per-layer curves (§3.5) are too blunt to separate them (shallow readings saturate post-shift, leaving that instrument blind exactly where the third account lives).

A trained default. The first account is the reading above: the covered task carries a trained default for unresolved cases, and the uncovered task does not.

Resolution at depth. A second, compatible candidate emerged from the depth data, though it is weaker than it first looked: fiction/real’s mid-deep stack (layers 3–18) settles clearly on the destination side in both directions, where tank aquarium→vehicle’s ends at the midpoint — so downstream layers in the covered task may act on a more settled reading than the site we measure. But the deepest band hovers near the midpoint in fiction→real too — the very direction where safeguard behavior must appear, and did (§3.2) — so depth-resolution cannot carry the explanation alone.

An early, surface-keyed trigger. A third candidate points the opposite way through the stack: the shallowest layers resolve the framing composition almost immediately and completely (§3.2), with the profile of surface-cue tracking — the minimal-pair discrimination that rules out cue-only tracking was run at the calibrated site, not at shallow layers — and a safeguard that reads early, from surface content and frame cues, would fire broadly whenever the alarming request is present and be suppressed only by a well-established fictional frame: the shape of our behavioral data, including its fragile 50% fiction-side floor. The fast, complete shallow response is itself specific to the fiction/real task — tank’s shallow layers respond later and, in one direction, only partially (§3.2) — which is what a trigger sculpted by safety post-training would look like, though ordinary register statistics learned in pretraining predict the same asymmetry and training provenance is unobservable here. Patching shallow versus mid-stack states during generation is the decisive test (§5).

We add one calibration so this section cannot overpromise: as a standalone monitor of individual readings, the frame reading is not yet usable (AUC 0.61 [0.43, 0.76], chance-compatible; Fig. 22). The band-level gradient is real; building a usable monitor — combining several sites, the mixed-context marker, and per-layer readings — is future work, not a claim.

A signal behavior does not appear to use. The mixed-context marker gives the dissociation its sharpest form: the model carries a persistent, systematic signal that its context is mixed — and behavior does not appear to use it at matched composition, at the one site and layer tested. This is the same structure reported in the hallucination literature, where models internally encode uncertainty or truthfulness that their generations do not respect. Here it appears in transition dynamics: an internal signal correlated with the conditions that produce the unresolved zone already exists inside the model, and behavior does not appear to use it — whether it could actually flag the zone is untested (the one monitor built, from the frame reading, is chance-compatible). The motivating case of §1 reportedly had the same structure at system scale — the provider’s moderation layer was flagging the user’s messages for self-harm risk in real time while assistance continued [CITE: filing/reporting]. That suggests a constructive direction for alignment work — not creating uncertainty signals, but testing whether behavior can be gated on the ones already there.

The human parallel. The signature of §3.2 — partial reanalysis that leaves a lingering trace of the initial misreading — is also the signature of human sentence processing in the “good-enough” tradition: after recovering from a garden-path sentence, comprehenders measurably retain components of the initial, incorrect interpretation. We make no mechanistic identification; we note that a language model trained on human text reproduces, at the activation level, the reanalysis profile humans show behaviorally — including the part where the old reading never fully leaves.

Garden path or pun. In Dynel’s terms our transition corpus is a slow-motion garden path, and most trajectories treat it as one: incongruity, then resolution toward the new reading. The dwelling runs are the interesting case — an *extended incongruity phase*, in which the model, asked

what the word means, answers like someone explaining a pun: both senses, held. Whether any of them is a true pun — held incongruity that never resolves — or only a long garden path is the persistent-versus-slow question our twenty-sentence horizon cannot decide, and the extended-tail experiment named in future work is its dedicated test.

Closing. Word-sense reinterpretation is the tractable laboratory instance of a much larger family — frames, personas, tasks, and safety postures all shift under accumulating context, and there is no obvious reason the metastable structure documented here is unique to word senses. The states are what the title calls them: unresolved. For now, so is the question of what a model should do while inside one.




Figure 22: The frame reading as a standalone monitor (exploratory, post-freeze): receiver operating characteristic of the calibrated-site reading as a predictor of safe-completion, AUC 0.61 [0.43, 0.76] (family-clustered interval), chance-compatible.

5 Limitations and future work

One model, stated plainly. Everything here is measured in one 20-billion-parameter mixture-of-experts model, with deterministic decoding. Each dynamics claim is made at one calibrated site and layer per task — single rows of a depth progression in which crossing timescales vary systematically (§3.2); no single layer is a sufficient readout of the model’s state, and behavior is produced downstream of all of them. Within the study, replication is internal: two task contrasts, three carriers plus a replicate, two measurement sites, twelve scene families per class per task. That inventory also records what did *not* replicate cleanly: the real→fictional remnant is suggestive only once reference uncertainty is propagated; the letter-site asymmetry rests on four runs per direction; 48% of fiction/real runs are indeterminate under per-run model selection; and our pre-stated prediction about the ordering of crossing times across layers held in one task and failed in the other. Cross-task patterns — the task with the wider endpoint separation also shows the larger remnant gaps; anchored versus distributed contrast structure — are $n = 2$ observations, hypotheses rather than findings.

The named open question. Whether the remnant — and the dwelling within the unresolved zone — are permanent or merely slower than twenty sentences: in the terms of §4, whether the dwelling is a pun rather than a long garden path — is decidable by extending the post-shift horizon (extended-tail runs) and by shifting evidence back to the original class ($A \rightarrow B \rightarrow A$ runs, measuring what remains of the second frame). Both are designed and costed; neither has run.

Further deferred work, in rough order of leverage:

- Replicating the fiction/real headline quantities at the ‘write’ site, where the framing contrast reads most strongly (the existing recordings suffice).
- Expanding the letter-site families.
- Calibration sets from a third, unrelated class, to identify what the accumulation offset contains.
- Monitor construction (multi-site, marker-augmented, depth-aware) against the chance-compatible single-site baseline reported above.
- Causal localization of where safeguard behavior reads in the stack: patching shallow versus mid-stack states during generation — the decisive test for §4’s three candidate accounts.
- A planned suite of harder context manipulations — colliding frames, mutating them mid-context, deliberately incoherent contexts — where genuine off-distribution excursions are most likely, our clean block-shifts being the tamest possible case.

The analyses reported here were frozen before drafting; the post-freeze additions are logged as such in Appendix B.

Back matter (assembly spec — materialized at final build)

Review-draft note: the sections below are the assembly specification; the final build replaces them with the referenced frozen content.

Acknowledgments

[Agreed AI-assistance line — exact text from Andrew/chat side, placeholder.] Personal acknowledgments: Andrew adds.

Related work (arXiv default: standalone section, drafted from handoff §6 anchors)

Probing validity: Alain & Bengio; Hewitt & Liang (control tasks — our D5 + label-shuffle audits presented as the selectivity answer); Belinkov survey. Directions & steering: Marks & Tegmark (also the mass-mean/diff-of-means probing precedent for §2.2’s axis choice); Ardit et al. (refusal direction — nearest cousin of the fiction/real behavior link). In-context dynamics: Xie et al. (ICL as implicit Bayesian inference — the standing theory our integrator null speaks to); Hendel; Todd (function vectors); Olsson et al. (induction heads). Safety: many-shot jailbreaking; Crescendo-style multi-turn escalation. Psycholinguistics: Christianson & Ferreira; Slattey et al.; van Schijndel & Linzen. Verbalized uncertainty & ambiguity (grounds the intro’s scoped claim): Kadavath et al.; Lin et al. (verbalized confidence); Xiong et al. (calibration of verbalized confidence); Liu et al. (ambiguity modeling). Internal encoding vs expression (grounds Discussion’s “represented but not consulted”): Azaria & Mitchell (internal state predicts truthfulness); Orgad et al. (models encode more than they express). Long-context artifacts: attention sinks (the “is the offset positional?” question; midpoint referencing makes our findings robust either way). Dynamics: Sussillo & Barak; Kelso; Rabinovich et al. Motivating case (intro + discussion): Raine v. OpenAI, Cal. Super. Ct. (S.F.), filed 2025-08-26; contemporaneous reporting (CNN 2025-08-26; NBC News; TechCrunch 2025-11-26 on OpenAI’s answer denying causation). Contested litigation — the paper characterizes only what the filing and reporting describe. Model provenance (the intro’s “trained safeguard”): gpt-oss model card, safety training documented.

Appendix A — Corrections affecting reported values

One line each for the corrections that changed numbers as printed: the real→fictional remnant interval widening to include zero under reference resampling; the trimmed within-stream means; the per-run crossing-time medians; the recency-parameter memory figures. The complete corrections record is in the repository.

Appendix B — QA and reproducibility

Regeneration audit summary (bit-identical axes; 23-script diff green; seeded bootstraps exact); 9/9 fixture suite; sign/boundary/join/shuffle audits; the four logged post-freeze additions (s12 response counts, 0/96; s11 monitor ROC, AUC 0.61 [0.43,0.76]; s13 per-layer collapse panels; s14 exploratory per-layer behavior association — both computed from frozen captures via the committed cache); repository pointer + regeneration instructions. Terminology map for repository readers: the paper’s *remnant* is the repository’s “residual”; the dwelling within the unresolved zone is the repository’s “park”; *accumulation offset* is “accumulation drift”; *tasks* are the repository’s “probes”/“probe arms”; corpus directories D3/D4/D5/D6 are the transition, no-shift, minimal-pair, and mixture-sweep corpora; the repository’s “K=1” label names an analysis convention for a routing track this study dropped.

Appendix C — Supplementary figures

figures/other/ (12): fr fit gallery, fr D6 loop, tank secondary heatmap, matched-k panel is MAIN (`fig_s9_behavior_matchedk`) — supplement gets: raw + midref heatmap generations (the instrument story’s “before” pictures), `prelexical_L4`, `occupancy_bands_L4`, `jumpiness`, `norm_vs_alignment`, `calibration_layers`, plus `fig_s11_monitor_roc`. `fig_s13_collapse_layers` is MAIN (cited in §3.2’s depth paragraph; committed script `s13_collapse_by_layer.py`).